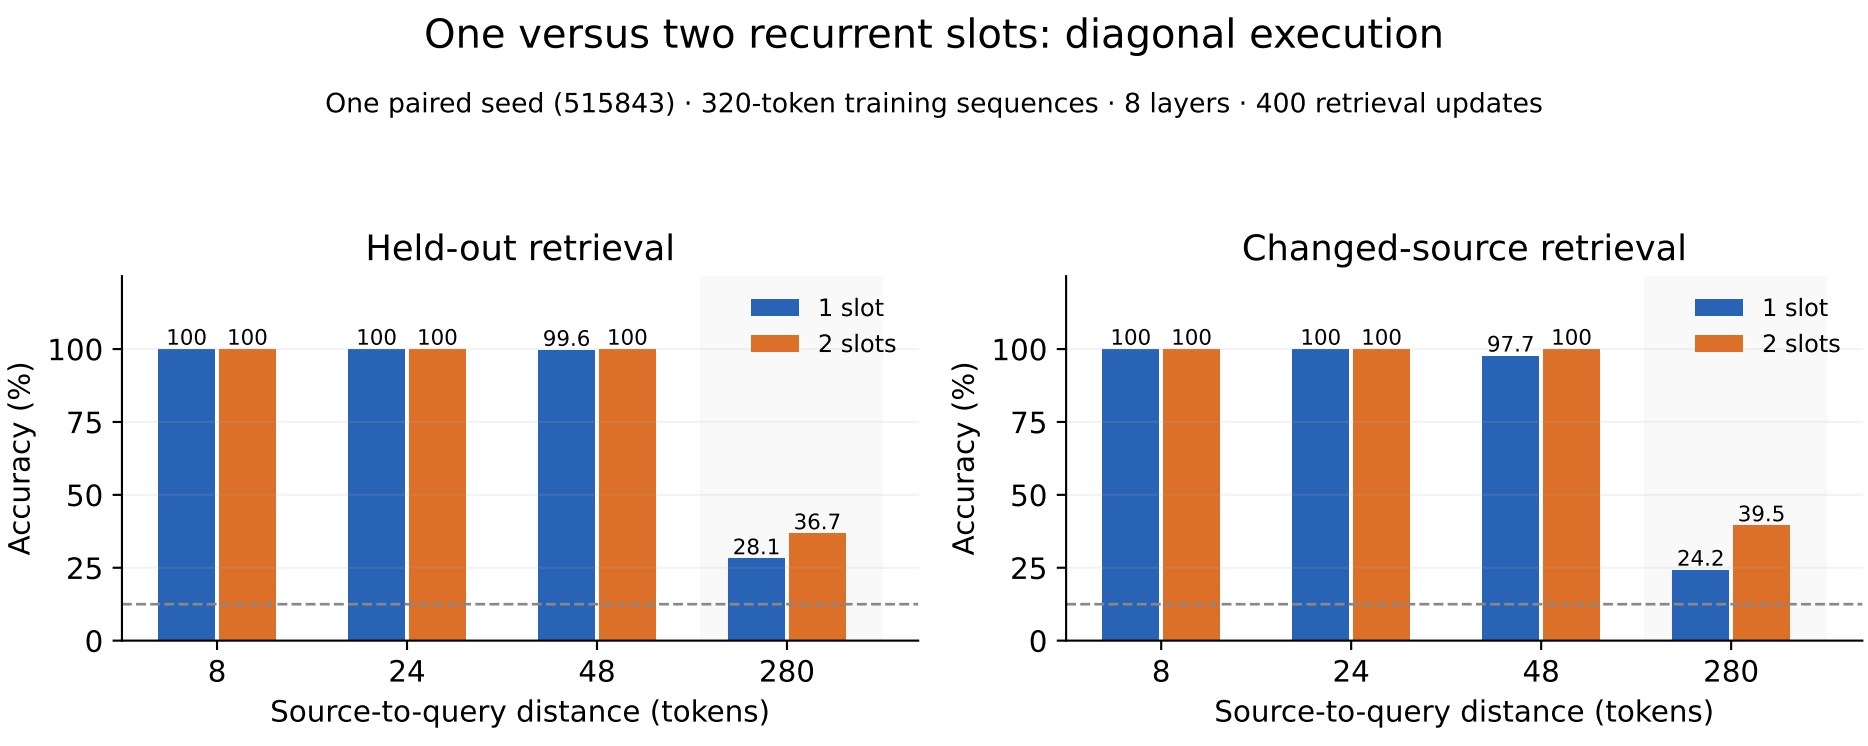
One versus two recurrent slots: diagonal execution
One paired seed (515843) · 320-token training sequences · 8 layers · 400 retrieval updates
Held-out retrieval
Changed-source retrieval
1 slot
2 slots
Accuracy (%)
0
25
50
75
100
100
100
100
100
99.6
100
28.1
36.7
97.7
24.2
39.5
8
24
48
280
Source-to-query distance (tokens)
Same initialization, sampled batches and labels. Dashed line: 12.5% chance.
Shaded group: beyond the baseline horizon. One seed does not establish reliability.