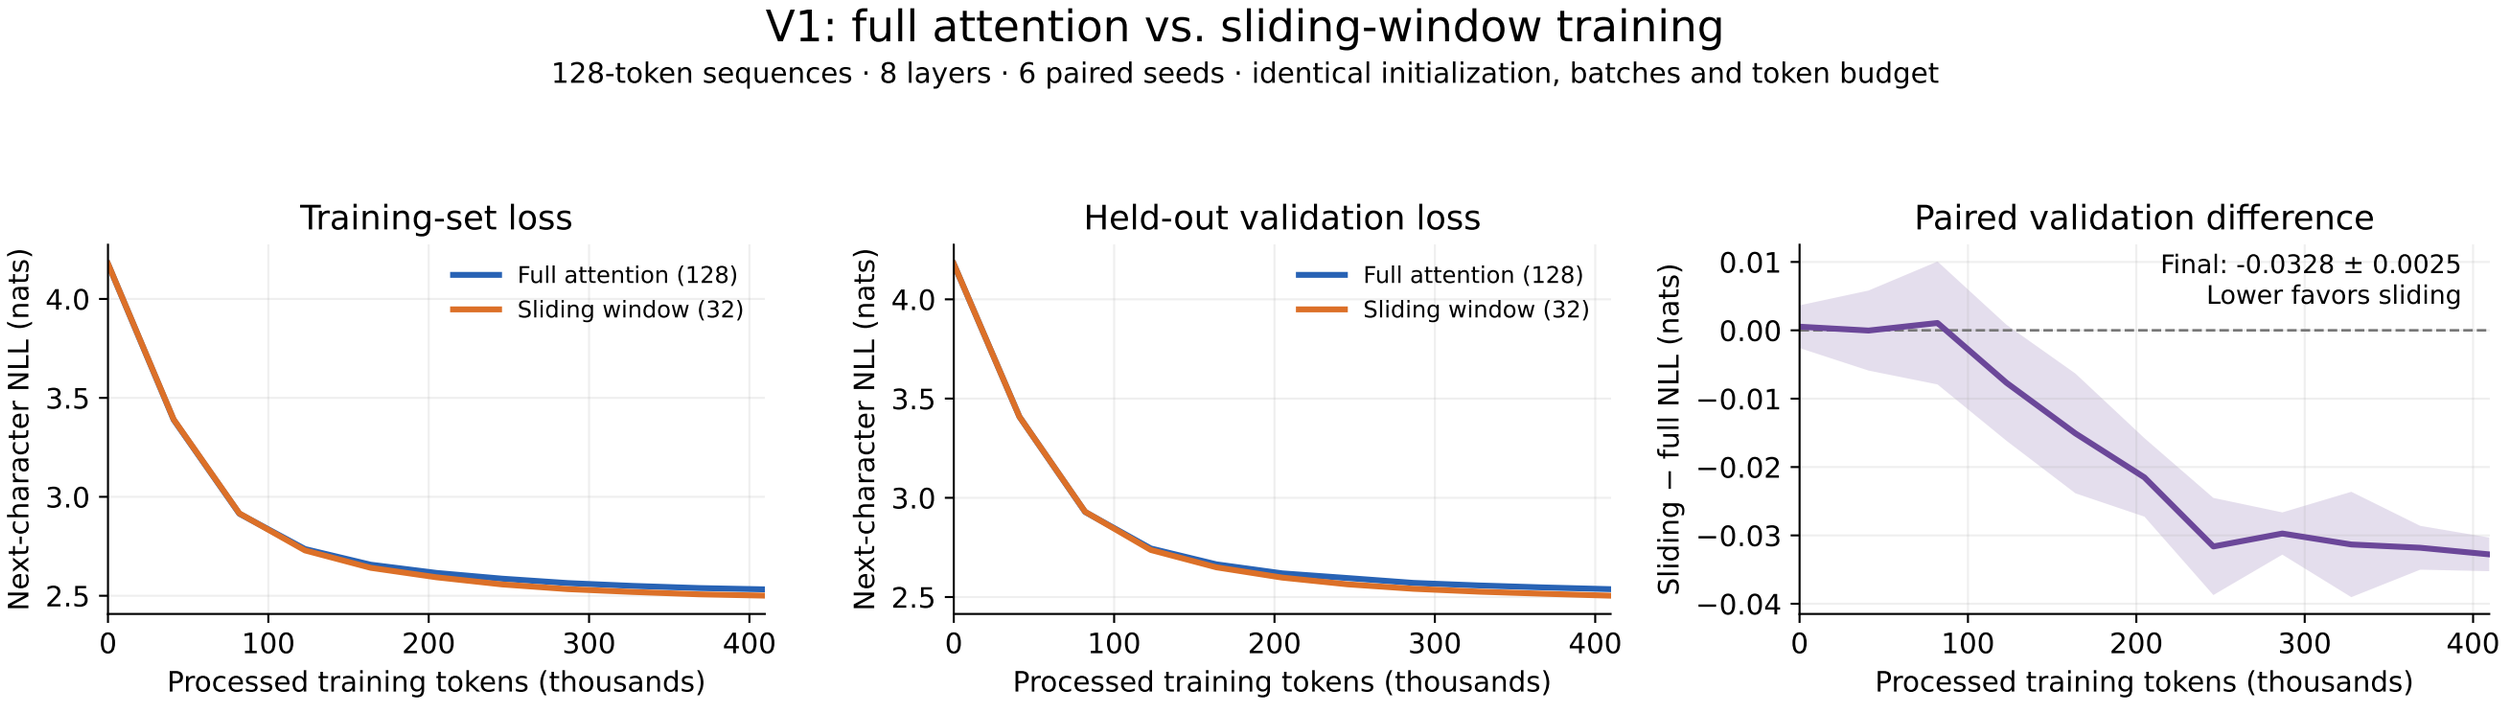
V1: full attention vs. sliding-window training
128-token sequences · 8 layers · 6 paired seeds · identical initialization, batches and token budget
Training-set loss
Held-out validation loss
Paired validation difference
Full attention (128)
Sliding window (32)
Next-character NLL (nats)
Sliding − full NLL (nats)
Processed training tokens (thousands)
Final: -0.0328 ± 0.0025
Lower favors sliding
Lines: six-seed means. Shading: ±1 sample SD across seeds. Fixed 512-sequence sets scored every 20 updates.
Inference KV storage: 381 KiB (full) vs. 93 KiB (sliding); training-memory savings are not measured.