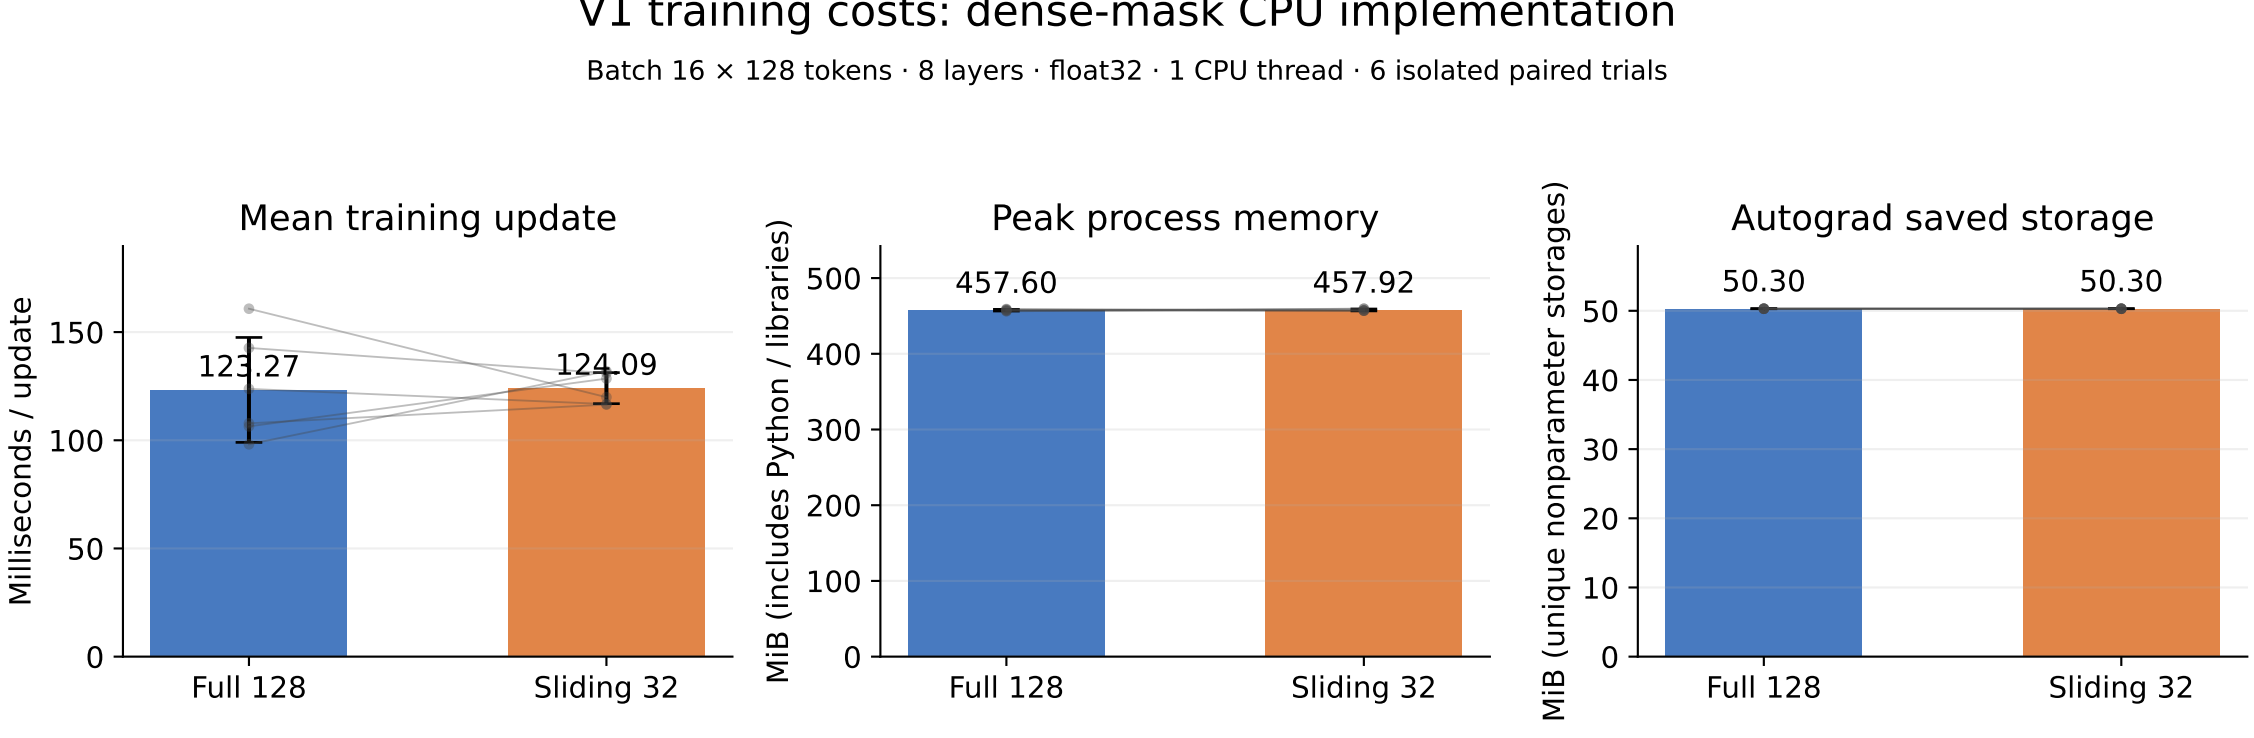
V1 training costs: dense-mask CPU implementation
Batch 16 × 128 tokens · 8 layers · float32 · 1 CPU thread · 6 isolated paired trials
Mean training update
Milliseconds / update
0
50
100
150
123.27
124.09
Full 128
Sliding 32
Peak process memory
MiB (includes Python / libraries)
0
100
200
300
400
500
457.60
457.92
Full 128
Sliding 32
Autograd saved storage
MiB (unique nonparameter storages)
0
10
20
30
40
50
50.30
50.30
Full 128
Sliding 32
Timing: 10 warmup + 50 measured forward/backward/clip/AdamW updates per trial. Error bars: sample SD.
CPU memory is distinct from inference KV storage; saved-storage counts are not total peak memory.